

# MEV Opportunity Identification and Execution Framework

## Introduction

Miner/Maximal Extractable Value (MEV) refers to profits that can be captured by reordering or inserting transactions in a blockchain block <sup>1</sup>. MEV bots actively scan blockchain transactions (especially on EVM chains like Ethereum) to exploit opportunities such as **arbitrage** (price discrepancies across exchanges), **sandwich attacks** (front-running and back-running a large trade), and **liquidation sniping** (seizing under-collateralized loan positions) <sup>2</sup> <sup>3</sup>. This report presents a detailed design for a Replit-hosted application that **identifies and executes MEV opportunities** across multiple EVM-compatible chains using real-time data from CoinGecko Pro and DefiLlama Pro APIs. The system will incorporate a robust backend (including a data aggregator and MEV bot engine), a real-time dashboard frontend, and secure wallet integration to execute profitable transactions. Key features include detection of sandwich, arbitrage, and liquidation opportunities, live profit calculation (net of gas costs), confidence scoring with execution timing, multi-chain support (Ethereum, Polygon, Arbitrum, etc.), and a unified API layer to handle data sourcing and rate limits. We also emphasize security best practices (for safeguarding keys and avoiding malicious traps) to ensure the framework is safe and reliable for live deployment on Replit.

## MEV Opportunity Types and Detection

### Sandwich Attacks

A **sandwich attack** occurs when a bot detects a pending large swap in the mempool and quickly executes two trades: one **front-running** (buying the asset before the large swap to drive its price up) and one **back-running** (selling the asset right after the victim's swap at the inflated price) <sup>3</sup>. The bot profits from the price slippage caused by the victim's large order. To detect sandwich opportunities, the system's **MEV bot** will monitor incoming mempool transactions (pending swaps on decentralized exchanges) for unusually large trades that are likely to move prices. This requires a mempool subscription (using an Ethereum node's JSON-RPC or a third-party mempool service) to catch relevant transactions in real time <sup>4</sup>. When a large trade is observed, the bot will simulate the price impact on the target Automated Market Maker (AMM) pool and calculate a potential profit for sandwiching. Key data for this includes the token's current price and liquidity (to estimate slippage) from DEX data and CoinGecko, as well as current gas prices to assess transaction cost. The bot assigns a **confidence score** based on factors like trade size vs pool liquidity and network congestion, and defines a tight **execution window** (usually the next block) during which the sandwich must be executed for profit. If the confidence is high and net profit (after gas) is positive, the bot will attempt the sandwich by submitting two transactions with optimized gas (priority fee) to ensure they bookend the victim's transaction.

### Arbitrage Opportunities

**Arbitrage** MEV involves exploiting price differences for the same asset across different exchanges or markets <sup>2</sup>. The framework will continuously fetch **real-time prices** of assets across multiple DEXs (and possibly CEXs) using CoinGecko's market data and tickers endpoints. For example, the CoinGecko API provides a *coin tickers* endpoint that lists a coin's prices on various exchanges (both centralized and

decentralized) <sup>5</sup>, which is ideal for detecting price discrepancies. The system will use this along with DefiLlama's DEX volume/liquidity data to identify pairs of exchanges where an asset's price differs beyond a threshold (e.g. Token X costs 5% more on Uniswap than on SushiSwap). When a profitable arbitrage loop is found (buy low on one exchange, sell high on another), the bot module will calculate the *estimated profit* considering trade sizes, pool liquidity, swap fees, and gas costs. DefiLlama's **token liquidity** endpoint (which provides available liquidity for swapping a given token on specific chains <sup>6</sup>) helps determine if the volume needed for the arbitrage is feasible without excessive slippage. The bot assigns a confidence score based on liquidity depth and volatility (e.g. more confidence if price gap is large and both exchanges have high liquidity), and determines the timing (usually execute ASAP in the current block). For execution, the bot could either craft a single transaction that performs both swaps (using a flash-loan or a smart contract if needed), or execute two back-to-back transactions. The framework will incorporate **flash loan** capabilities or multi-call transactions if necessary to allow arbitrage without initial capital, although this requires additional smart contract logic (this can be a future expansion).

## Liquidation Sniping

**Liquidation** opportunities arise on lending protocols (like Aave, Compound, etc.) when a borrower's collateral value falls below the required threshold, allowing anyone to repay their debt and seize their collateral at a discount. MEV bots scan for **undercollateralized loan positions** and attempt to liquidate them before others do <sup>7</sup>. Our system will track protocol health factors using DefiLlama and on-chain data. DefiLlama's Pro API can provide protocol TVLs and possibly user deposit/borrow data; for instance, using the **inflows/outflows** endpoint for lending protocols to see rapid outflows (which might indicate collateral being withdrawn or prices dropping) <sup>8</sup>. Additionally, the bot may integrate directly with protocol subgraphs or use DefiLlama's "Liquidations" dashboard data (if available) to get a list of at-risk positions. When a potential liquidation is detected (e.g., an account on Aave is below collateral ratio), the bot will calculate the profit by comparing the collateral seized vs. debt repaid (taking into account the protocol's liquidation bonus). It will also factor in gas cost and the likelihood of success (confidence could be based on how far below the threshold the position is – deeper means more time to act). The execution involves calling the protocol's **liquidation function** via the user's wallet. These transactions are time-sensitive; the bot might use a high gas price or Flashbots bundle to ensure the liquidation transaction is mined promptly before competitors <sup>9</sup>. Liquidation sniping requires monitoring multiple protocols and chains; the framework will be designed to easily plug in new protocol adapters (e.g., a module for Aave v3 on Polygon, Compound on Ethereum, etc.) that fetch health factor data and perform liquidation calls.

## Data Sources and APIs

To support the above detection strategies, the system aggregates **real-time data** from the CoinGecko Pro API and the DefiLlama Pro API (and a few other sources as needed, such as gas or mempool data). All external API calls are funneled through a **unified data service** in our backend to handle authentication (API keys) and rate limiting centrally. Below, we outline the key data requirements and how each API is used to fulfill them:

- **Token Prices (Real-Time)** – We use CoinGecko's price endpoints to fetch current token prices in USD (or ETH) at high frequency. Specifically, the `/simple/price` endpoint allows querying live prices for multiple coins by their IDs <sup>10</sup>, which is useful for getting a quick price feed for assets involved in potential MEV trades. We also use `/simple/token_price/{platform}` when needed to get the price of tokens by contract address on a given chain (for assets that may not have a CoinGecko ID, e.g., newly launched tokens on Uniswap). This price data is crucial for arbitrage calculations (to compare prices across venues) and for valuing collateral in liquidation

scenarios. Price updates will be fetched frequently (e.g., every few seconds) but carefully throttled to respect rate limits (hence using a unified service that can cache results for a short time).

- **Exchange and Liquidity Data** – For arbitrage detection, knowing prices on specific DEXs and their liquidity is important. CoinGecko's **Coin Tickers** endpoint (`/coins/{id}/tickers`) provides the prices and volumes of a coin on various exchanges (both DEX and CEX) <sup>5</sup>. This can highlight, for example, that token ABC is trading at \$10 on Uniswap but \$10.50 on SushiSwap, indicating an arbitrage gap. In addition, DefiLlama's data on DEX volumes and liquidity will be utilized. DefiLlama Pro offers an overview of DEX volumes (e.g., via an endpoint listing all DEXs with their 24h volumes and historical data) <sup>11</sup>, and these can be filtered by chain <sup>12</sup> to find which DEXs on a given chain are most active. We will call endpoints like `/api/overview/dexs` (all DEXs summary) and `/api/overview/dexs/{chain}` (DEXs on a specific chain) to get insights into where trading activity is high. This helps focus the bot's attention on the most liquid venues where large trades (and thus MEV opportunities) are more likely. Furthermore, DefiLlama's **Token Liquidity** feature provides the available liquidity for swapping between tokens on various chains <sup>6</sup>, which we use to estimate slippage and feasible trade sizes for arbitrage or sandwich attacks.
- **Protocol Data (TVL, Flows, Active Users)** – To monitor the DeFi ecosystem for anomalies (which could hint at MEV chances), we leverage DefiLlama's comprehensive protocol statistics. For instance, the **TVL (Total Value Locked)** for protocols and chains gives a macro view of where capital is moving. We can retrieve TVL by chain or protocol using DefiLlama (the Pro API includes endpoints under the TVL category, such as total assets on all chains <sup>13</sup> or specific protocol TVLs). In particular, if a protocol experiences a sudden drop in TVL, it might suggest a large withdrawal or liquidation event. DefiLlama's `/api/inflows/{protocol}/{timestamp}` endpoint provides **inflows and outflows** for a protocol at a given time <sup>14</sup>, which we can use to detect unusual fund movements (e.g., a huge outflow from a lending platform could indicate multiple liquidations or a fleeing whale). Additionally, the **Active Users** data from DefiLlama (via `/api/activeUsers` and `/api/userData/{type}/{protocolId}`) gives the number of active users or addresses interacting with a chain or protocol <sup>15</sup>. A spike in active users on a DEX or bridge could correlate with arbitrage opportunities (many users trading a token might create price volatility), whereas active user drops on a lending protocol might accompany liquidation events. These data points will be polled periodically and displayed on the dashboard to provide context (e.g., showing user activity trends, which can be correlated with MEV capture events).
- **Gas Prices and Fees** – An essential part of profit calculation is the gas fee estimation. Neither CoinGecko nor DefiLlama directly provides gas price data, so our framework will integrate a **gas oracle** or use the blockchain RPC to get current gas prices. For Ethereum, we might call an endpoint like `eth_gasPrice` via an Infura or Alchemy node, and also consider priority fee (EIP-1559) for timely inclusion. We could also use third-party APIs or CoinGecko's global data if available (CoinGecko has a *Global DeFi* data endpoint <sup>16</sup> that might include info like gas, though not explicitly stated). The system will maintain an updated gas price (in Gwei) and use chain-specific gas costs (for Polygon, Arbitrum, etc., since they have different fee regimes) to compute the **net profit** of each opportunity. Additionally, for more accuracy, the bot can use the RPC's `eth_estimateGas` for the specific transaction to be executed, ensuring we account for actual gas limit and cost of that trade. All this will be unified under our backend's API so the frontend can simply fetch "current gas fee" and "estimated cost" for a given trade when displaying opportunities.

- **Other Data** – We will also utilize specialized data as needed. For example, **historical price charts** from CoinGecko (using endpoints like `/coins/{id}/market_chart`) can be used to display trends on the dashboard or to calculate volatility (which might feed into confidence scoring for an opportunity). DefiLlama's **Yields** endpoints (like `/yields/pools` for latest APYs of pools <sup>17</sup>) can help spot unusually high yields that might indicate an inefficiency (perhaps arbitrage between lending rates, though that's a more complex MEV). DefiLlama also tracks **stablecoin dominance** on chains (`/stablecoins/stablecoindominance/{chain}`) <sup>18</sup> – a sudden shift in stablecoin liquidity on a chain might affect DEX pricing and thus arbitrage. While not all these data are directly used to trigger a trade, they provide a richer context for the user and the bot to understand the market state. We incorporate them into the dashboard metrics (for example, showing TVL or volume changes over 24h alongside the profit captured in that period).

All API calls are routed through a **Unified API Service** layer in the backend. This service holds the API keys (for CoinGecko Pro and DefiLlama Pro) securely and handles tasks like **rate limiting**, caching, and combining data. For instance, if multiple frontend components need price data and TVL data, the service can parallelize the calls to CoinGecko and DefiLlama, then return a single combined response. It also ensures we don't exceed the allowed call frequency of each API by queueing requests or using cached responses (e.g., price data could be cached for a few seconds if the APIs allow, to avoid spamming requests). This design abstracts the data layer away from the rest of the app – if we switch to a different data provider or add another source (like a subgraph for liquidation data), we can integrate it into this service without changing the frontend or bot logic.

## Architecture Overview

The system is composed of three main components – the **Frontend UI**, the **Backend Server**, and the **MEV Bot engine** – all orchestrated to work together on Replit. The high-level architecture is illustrated in the diagram below, which shows data flows between components and external services:

*Figure: System architecture for the MEV app, showing the browser-based Frontend, Backend/API layer, MEV Bot engine, and external integrations with CoinGecko, DefiLlama, and blockchain nodes.*

### Frontend (Dashboard UI)

The frontend is a web application (served by the Replit app) that provides users with a **dashboard** to visualize MEV opportunities and system performance. It will likely be implemented using a JavaScript framework (Replit supports Node.js easily, so we might use **React** or even a simple static HTML/JS with libraries like Chart.js for charts). The UI is designed for clarity and real-time interactivity:

- **Opportunity Feed:** A table or list displaying current detected opportunities (sandwich, arbitrage, liquidation) across supported chains. Each entry shows details like the asset(s) involved, the type of opportunity, the estimated profit (and net profit after gas), the confidence score, and a countdown or time window for execution. For example, *"Arbitrage: Buy ETH on Uniswap (price \ \$1,800) and sell on SushiSwap (price \ \$1,820) – Profit ~\ \$20/ETH, Net ~\ \$15 after gas, Confidence: High, Expires in ~15s."*
- **Real-Time Charts and Metrics:** The dashboard will have panels for key metrics updated live. This includes a **24h MEV capture** chart (showing how much profit the bot has made in the last 24 hours, possibly broken down by strategy type), **Average Profit per Trade**, **Win Rate** (percentage of opportunities successfully executed vs attempted), and other stats. We will also show

correlations, for instance a chart or matrix that correlates certain variables: e.g., *profit vs gas price* (to see if high gas reduces profitability), or *number of opportunities vs market volatility*. We might leverage a charting library to draw these in real-time. Data for these come from the backend: the bot will log each executed trade outcome, and the backend aggregates those logs for 24h sum, averages, etc., exposing an endpoint for the frontend to fetch these stats.

- **Controls and Wallet Integration:** The UI will allow the user to connect their wallet or provide a key for executing trades. We plan to integrate **MetaMask** for a seamless and secure user experience. Using MetaMask's web3 integration, the frontend can prompt the user to connect and sign messages or transactions when needed <sup>19</sup>. MetaMask is beneficial because it manages the private keys securely in the user's browser and confirms transactions with the user <sup>20</sup>, meaning our app never sees the raw private key (enhancing security). The interface will have a "Connect Wallet" button (for MetaMask or WalletConnect) and alternatively a section to enter a private key (for advanced users who run headless or are okay storing a key on the backend – this would be clearly marked as higher risk). Once connected, the user can choose settings like *auto-execution vs manual*. In manual mode, when an opportunity is detected, the UI could present a "Execute" button that, when clicked, uses the connected wallet to send the transaction (or instructs the backend to do so). In auto-mode (for trusted keys), the bot will execute without user intervention. The frontend also includes configuration options for the user (e.g., which chains or strategies to enable, risk thresholds, etc.), and may display system status (e.g., "Connected to Ethereum + Polygon. Gas: 30 gwei. Bot Active.").
- **Real-Time Updates:** To keep the UI in sync with the backend, we will use either WebSocket connections or periodic polling (depending on complexity and Replit support). A WebSocket or similar push mechanism can send new opportunity data to the frontend instantly (e.g., the moment the bot finds an arbitrage, the server pushes it to the UI so it appears without user refreshing). Alternatively, the frontend can poll an endpoint every few seconds for updates. Given the need for very timely updates (especially for fast MEV opportunities), a push mechanism is preferred. The Replit backend can maintain a WebSocket channel or use Server-Sent Events to stream updates. This way, the dashboard feels live, with tables updating and charts animating as new data comes in.

## Backend (API & Data Aggregation Layer)

The backend is the heart of the application, managing data gathering, decision logic, and serving API endpoints to the frontend. It can be implemented as a Node.js server (using **Express** or a similar framework) or Python (using **Flask/FastAPI**), depending on familiarity. On Replit, Node.js might integrate well with front-end JS, but Python could be used if leveraging certain ML or math libraries for scoring. Key responsibilities of the backend include:

- **Unified API Service:** As described earlier, this service handles all communication with **CoinGecko** and **DefiLlama** (and any other external APIs). It acts as a facade that the rest of the backend (and the frontend) uses to request data. For example, when the bot needs the latest prices and liquidity for token XYZ, it calls a function in this service, which in turn may call `CoinGecko /simple/price` for price and `DefiLlama /api/tokenProtocols/xyz` for liquidity across protocols. The unified service will merge and normalize these results. It also keeps track of how many calls have been made to each API and ensures we stay within rate limits (for instance, CoinGecko Pro might allow X calls/minute depending on plan; the service will queue calls if we approach that). If multiple parts of the system request the same data in quick succession, caching is used – e.g., store the last price fetch for a token for a few seconds. This layer also uniformly handles authentication (API keys) so that those keys are never exposed

outside the backend. In terms of implementation, this could be a set of utility modules or an internal API. We might use libraries like **Axios** (for Node.js) or **Requests** (for Python) to call external APIs, possibly with retry logic and back-off to handle any errors gracefully. The unified service might also incorporate a small database or in-memory store for frequently needed data (for example, caching the list of CoinGecko IDs or protocol slugs from DefiLlama to avoid repeated metadata calls).

- **MEV Detection Engine:** This is the logic that processes data and identifies opportunities. It runs as part of the backend (potentially in separate threads or an async loop). For each category of MEV, the engine uses data from the unified service and direct blockchain monitoring:
- **Arbitrage:** The engine periodically pulls price quotes for key trading pairs from multiple sources (via CoinGecko or directly via on-chain DEX price oracles). We could maintain a list of popular tokens and their addresses on each DEX. The engine computes price ratios; if any ratio deviates significantly ( $>$  threshold), it flags an arbitrage candidate. It then checks DefiLlama liquidity data to ensure volume can be traded. If conditions look good, it creates an **Opportunity object** containing all details (trade path, expected profit, etc.).
- **Sandwich:** The engine listens to mempool (which we achieve by connecting to an Ethereum node's WebSocket or using a service like Blocknative). When a pending swap is detected, the engine runs a quick simulation (it could use a local fork of the chain state or a library that simulates AMM swap impact given pool reserves). If a profitable sandwich is possible, it crafts the strategy (how much to buy in front-run, etc.) and encapsulates in an Opportunity object. Because these need immediate action, the engine will likely directly invoke the execution module (see below) rather than just wait for a schedule.
- **Liquidation:** The engine queries lending protocol data periodically. This might be via DefiLlama (if they offer a "liquidations pending" feed or at least health factors) or via each protocol's API/subgraph. It calculates health ratios and identifies any accounts below threshold. For each such account, it forms an Opportunity object with details (which protocol, how to execute the liquidation, est. profit).

This detection engine can be structured in a modular way: separate sub-modules or classes for each opportunity type. They can run on timers or event triggers concurrently. For instance, an **ArbitrageScanner** running every 5 seconds, a **MempoolWatcher** running continuously for sandwiches, and a **LiquidationChecker** running every 10 seconds. The output of each is standardized (so the rest of the system can handle an opportunity uniformly whether it's arb, sandwich, or liquidation).

- **Profit & Risk Calculation:** Within each opportunity identified, the backend calculates the expected profit. This includes querying the **live gas price** from an RPC (for Ethereum, possibly adjusting for priority tip). We may also use CoinGecko's gas metrics if available, but likely direct RPC is best for accuracy. The system will incorporate a gas *estimator* – for example, using the `eth_estimateGas` for a transaction or having predefined gas costs for certain actions (like a Uniswap swap costs ~120k gas). The profit calculation will subtract `estimated_gas * gas_price` (converted to ETH or token value) from the gross profit. It will also consider protocol fees (like 0.3% DEX fee) and, in case of flash loan usage, any fees for that. In addition, we assign a **confidence score** to each opportunity. This can be a heuristic formula taking into account factors such as:
- Price disparity magnitude (bigger arbitrage gap yields higher confidence),
- Liquidity buffer (if after our trade the price will equalize or still leave profit, etc.),
- Time sensitivity (opportunities that are expiring very soon might be lower confidence if we're not sure to get in the same block),

- Past success rate of similar opportunities (the system can learn or at least record outcomes: e.g., sandwich attacks on very volatile tokens may fail often due to unpredictability, so confidence could be lower).

The confidence score (say 0 to 100 or Low/Med/High) helps decide if auto-execution should proceed. For example, the user might configure “only auto-execute high-confidence opportunities”.

- **Backend API Endpoints:** Besides internal logic, the backend exposes RESTful endpoints to the frontend. Key endpoints include:

- `/api/opportunities` – returns a list of current opportunities (with details like type, profit, etc.) for the dashboard.
- `/api/stats` – returns aggregated stats (24h profit, win rate, etc.) for the dashboard metrics.
- `/api/config` – perhaps to get settings or supported chains/tokens lists for the frontend.
- `/api/execute` – (if using manual execution) to trigger the bot to execute a specific opportunity, possibly taking an opp ID or details (in auto mode, this isn’t used).
- These endpoints will secure any sensitive actions (e.g., executing trades) with an auth if needed (for instance, a simple token or ensure it’s called from the frontend only – Replit apps might run both front and back on the same origin, so XSRF protection etc., should be considered).

## MEV Bot Engine (Transaction Executor)

While we have described the detection part of the “bot,” the **execution component** is critical and can be thought of as a custom bot service. This could be integrated in the backend process or run as a separate worker (to allow parallelism). The bot engine is responsible for taking an Opportunity (from the detection engine or a user-trigger), and **carrying out the necessary blockchain transactions** to realize the profit. Key aspects of the bot design:

- **Wallet Integration:** The bot uses the user’s wallet credentials to sign and send transactions. If the user connected via MetaMask and is manually confirming, the bot might simply surface the transaction details to the frontend for the user to confirm (e.g., using Web3.js to invoke MetaMask’s `ethereum.request({ method: 'eth_sendTransaction', params: [...] })`). However, for fully automated execution, the bot will use a **private key** stored in the backend (provided by user and kept in an environment variable or secure storage on Replit). We will likely use a library like **Ethers.js** (in Node) or **Web3.py** (in Python) to handle transaction signing and building <sup>21</sup>. These libraries allow creating a wallet object from a private key and easily constructing transactions to call smart contracts or send funds. For instance, Ethers.js provides a `Wallet` class that can connect to a provider (Infura/Alchemy RPC) and send transactions; Web3.py similarly can load a private key into a local account for signing. Using these libraries, the bot can programmatically set the `nonce`, `gasPrice` (and `maxPriorityFee`), `to`, `data`, `value` fields of transactions. We will integrate chain-specific settings (for example, on EIP-1559 chains like Ethereum, set `maxFeePerGas` and `maxPriorityFeePerGas`).

- **Transaction Construction:** Depending on the type of opportunity, the bot may need to build different transactions:

- *Arbitrage:* likely two swaps (or one if using a smart contract/flash loan). Without a flash loan, if the user already holds one asset, it could do one swap trade and then another. With flash loans (advanced scenario), the bot would call a flash loan provider within a smart contract that then executes both trades and returns the loan. In our design, initially we assume simpler arbitrage

where the user's base capital is used: e.g., use some ETH to buy Token on Exchange A, then sell Token on Exchange B for ETH, ending with more ETH than started. We will have pre-written contract ABIs (Application Binary Interfaces) for the DEX routers (UniswapV2, UniswapV3, Sushi, etc.) to encode the swap function calls. For example, UniswapV2's

`swapExactTokensForTokens` function can be encoded with Web3/Ethers easily given the amounts and path. The bot will choose the optimal path (possibly just direct if between two tokens, or a route via common pool if needed).

- **Sandwich:** the bot needs to send two transactions: a buy and a sell. The challenge is to get them in the right order around the victim's tx. Here we will almost certainly want to use **Flashbots** or private transaction submission. Flashbots is a service where you can send a bundle of transactions directly to miners/validators so they are executed in a block without going through the public mempool <sup>22</sup>. Our bot can integrate with Flashbots by using their provider (they have an ethers provider for sending bundles). The bot would package the two sandwich transactions (and optionally the target tx if we can get it from mempool) and submit as a bundle with a bribe to the miner. This greatly increases success rate and avoids **gas wars** with other bots, but does require connecting to the Flashbots RPC endpoint. If Flashbots integration is too advanced for initial deployment, the fallback is to send the front-run tx with a very high gas fee to outbid others, then the back-run tx with slightly lower but still high fee, hoping they bookend. The bot will utilize the `maxPriorityFee` to try to get priority in the block for these.
- **Liquidation:** the bot will call the protocol's liquidation function. For example, Aave's `liquidationCall()` or Compound's `liquidateBorrow()`. We'll use the ABI of those contracts and call via our wallet. Typically these functions require specifying the borrower address and possibly the asset to seize. The bot sets those parameters based on the identified opportunity. We also must ensure the bot (user) has enough capital to repay the debt (or again use flash loan – many liquidations by bots do use flash loans to get capital for payoff). Our design can incorporate flash loans similarly to arbitrage if needed, but initially, the user might provide some capital for liquidations (or we limit to scenarios where the user's wallet has requisite funds). After constructing the transaction, the bot sends it to the network with a gas price that ensures quick mining.
- **Security & Monitoring:** The bot will monitor the status of its transactions. Using the web3 library, it can listen for confirmations or use the RPC to get transaction receipts. If a transaction fails (reverts or runs out of gas), the bot logs it (and the frontend can display a "failure" event, affecting the win rate metric). For critical, time-sensitive sequences like sandwiches, if the first transaction fails or is delayed (e.g., gets stuck behind something), the bot might drop/cancel the strategy (e.g., not execute the second leg to avoid unintended losses). The bot also employs some safety checks: for instance, just before sending a trade, it might double-check the latest price to ensure the opportunity still holds (markets move quickly, so something that was profitable 5 seconds ago might no longer be). If conditions have changed unfavorably, it can abort before spending gas on a doomed trade.
- **Multi-Chain Operation:** Our bot is designed to handle multiple chains. This means it may connect to multiple RPC endpoints (one for Ethereum, one for Polygon, etc.) and have multiple wallet instances (usually the same private key can be used on different EVM chains, as the addresses are the same, but the bot needs to track nonces per chain and have provider connections for each). We will structure the code such that chain-specific parameters (RPC URL, chain ID, contract addresses for DEXs) are configurable. The bot might run a separate loop for each chain's opportunities, or a unified loop that goes through chains in sequence. To ensure performance, concurrent processing is ideal – e.g., using Node's asynchronous capabilities or Python's `asyncio` to have tasks for each chain in parallel. The user can enable/disable certain



chains via configuration. For example, if only Ethereum and Polygon are turned on, the bot will focus on those and use the respective data (CoinGecko's price endpoint can return prices across networks since it identifies coins by a unique ID regardless of chain, but for chain-specific liquidity we use DefiLlama's chain filters). The modular design also means adding a new chain (say Binance Smart Chain) is as simple as adding its config: list of DEXs/routers, RPC endpoint, chain ID, etc., and perhaps updating coin address mappings for that chain.

- **Logging and Analytics:** The bot will log all identified opportunities and actions taken. These logs feed the dashboard metrics (success/fail, profit, etc.). They also are valuable for debugging and improving the strategies. Over time, the analytics can reveal patterns (for example, perhaps most sandwich attempts on a certain DEX fail due to competition – then the user might tweak gas strategy or avoid that DEX). Storing logs can be done simply in memory (for last 24h for instance) or persisted in a lightweight database or file on Replit if needed (though Replit's filesystem might reset on updates unless we use the persistent storage).

## Key API Endpoints for Data

The following tables summarize important API endpoints from CoinGecko and DefiLlama that the framework will utilize, along with their purpose in the system.

**CoinGecko API (Pro)** – used for token prices, market data, and exchange info:

Endpoint	Description	Usage in System
<code>/simple/price</code>	Returns current price of one or more specified coins (by CoinGecko ID) in various currencies <sup>10</sup> .	Primary source of live token prices (USD/ETH) for profit calculations and display. The bot calls this frequently to update prices of assets involved in opportunities.
<code>/simple/token_price/{platform_id}</code>	Returns price for tokens by contract address on a given blockchain platform (e.g., Ethereum) – part of the Simple API.	Used for tokens not tracked by a CG ID or to get price of a token on a specific chain. For example, price of a newly launched token by providing its Ethereum contract address. Ensures we can price any asset encountered.
<code>/coins/markets</code>	Retrieves market data for a list of coins, including price, market cap, 24h volume, etc. <sup>23</sup> .	Used to fetch a broad set of top coins and their 24h volumes. This helps identify which tokens are heavily traded (potentially more MEV opportunities). Also can be used to populate a list of tokens of interest dynamically.
<code>/coins/{id}/tickers</code>	Lists all trading pairs (tickers) for a coin on various exchanges (CEX & DEX) with prices and volumes <sup>5</sup> .	Critical for arbitrage detection – by getting the price of the coin on DEX A vs DEX B vs CEX. The backend can parse this to find price discrepancies (e.g., a DEX price vs another). We use it periodically for assets likely to have cross-market arbitrage.

Endpoint	Description	Usage in System
(Dex API) <code>/dex/pools</code> and <code>/dex/trades</code> (Beta)	CoinGecko's on-chain DEX API (currently Beta) has endpoints for pool data and recent trades <sup>24</sup> <sup>25</sup> . For example, trending pools or recent trades by token.	These can augment our detection: e.g., using <code>past 24 hour trades by token</code> <sup>26</sup> to identify large trades (possible sandwich targets) or a sudden volume spike. Trending pools might flag where big liquidity changes are happening (could be arbitrage targets).

**DefiLlama Pro API** – used for DeFi protocol analytics, liquidity, and volume data:

Endpoint	Description	Usage in System
<code>/api/chainAssets</code>	Returns the assets/totals for all chains (likely total TVL or asset holdings per chain) <sup>13</sup> .	Used for a high-level view of capital distribution. The app can display total TVL per chain and detect if capital flight from one chain to another is happening (which might create cross-chain arbitrage opportunities via bridges, although bridging arb is advanced).
<code>/api/tokenProtocols/{symbol}</code>	Lists the amount of a given token across all protocols ("Token Usage" data) <sup>27</sup> . For example, how much USDT is in various protocols.	Informs us where major liquidity pools are for that token. If an arbitrage involves USDT, this tells us which protocols (DEXes, lending, etc.) hold most USDT – indicating where price impact might be minimal or where big moves could occur.
<code>/api/inflows/{protocol}/{timestamp}</code>	Gives the inflow and outflow amounts for a protocol at a given time (timestamp) <sup>14</sup> .	The bot can query this for lending protocols or DEXs to see if, say, a huge outflow happened in the last hour (which could signal a large whale withdrawal or position close). This can trigger a check for related opportunities (e.g., after a big outflow, maybe a price deviation or a series of liquidations occurred).
<code>/api/activeUsers</code> and <code>/api/userData/{type}/{protocolId}</code>	Provides active user stats on chains/protocols <sup>15</sup> and user metrics of a specific type for a protocol.	We use active user counts as an indicator of activity surges. If a particular DEX suddenly has triple the usual active users in the past hour, it might mean a volatile trading event (potential arbitrage or sandwiches). These stats feed into our dashboard and possibly into opportunity scoring (higher activity might mean more competition, etc.).

Endpoint	Description	Usage in System
<code>/yields/pools</code>	Returns latest data for all yield farming pools, including APYs, TVLs, etc. <sup>28</sup> .	While not directly an MEV opportunity, abnormal yield values can imply something interesting (e.g., one pool's APY spikes due to an imbalance, which could be arbitrated by depositing/withdrawing). The system could flag extreme outliers in yield as potential opportunities (though this is more speculative). We also display some yield info in the dashboard for user insight.
<code>/api/overview/dexs</code> and <code>/api/overview/dexs/{chain}</code>	Lists all DEXes with summaries of their volumes and historical volume data <sup>11</sup> <sup>12</sup> .	The bot uses this to focus on the top DEXs per chain (since that's where MEV opportunities cluster). Also, tracking volume trends: if a smaller DEX's volume suddenly jumps, it could mean a large trade took place there (possibly leading to a price dislocation ripe for arbitrage). Volume history can also be shown on the dashboard.
<code>/stablecoins/stablecoindominance/{chain}</code>	Gives the stablecoin dominance and largest stablecoin info on a chain <sup>18</sup> .	This is used to understand the liquidity landscape of a chain – e.g., if USDC dominates Ethereum, many arbitrages might involve USDC as one side of trades. It helps in choosing base assets for scanning. Also, a change in stablecoin dominance might hint at news (like depeg or large movement) which itself can create arbitrage (as seen during stablecoin depegs).

*Table Note:* All DefiLlama Pro API calls require our API key and are made through the unified service, which handles the base URL and authentication. Similarly, CoinGecko Pro endpoints require our key and have higher rate limits which we utilize fully but carefully. We ensured above endpoints are included based on available documentation (DefiLlama's API offers many more endpoints like hacks, raises, etc., which are outside our current scope).

## Wallet Integration and Transaction Execution

Executing trades securely is a paramount part of this framework. We offer two modes of wallet integration: **MetaMask connection** for user-confirmed transactions and **direct private key** integration for automated execution. Both approaches are facilitated through well-established web3 libraries to interact with the blockchain.

- **MetaMask (Browser Wallet) Integration:** For users who prefer security and control, the frontend can integrate with MetaMask. MetaMask is a widely used browser wallet that allows dApps to request transaction signing from the user <sup>19</sup> . By using the MetaMask API (available via the injected `window.ethereum` provider), our app can connect to the user's wallet in one click.

Under the hood, we utilize a web3 library (such as **Web3.js** or **ethers.js**) to interface with MetaMask <sup>21</sup>. When an opportunity is to be executed in manual mode, the backend can send the transaction details to the frontend, which then calls `ethereum.request({ method: 'eth_sendTransaction', params: [txParams] })`. The user will see a MetaMask popup with the transaction details (to, value, gas, data) and can confirm or reject. This process ensures the private key never leaves the MetaMask extension, as it signs the transaction internally <sup>20</sup>. MetaMask manages the security of the keys and also provides network management – the user can switch to the required chain (our app can prompt network change if, say, an opportunity is on Polygon and the user is currently on Ethereum). One challenge is timing: MEV opportunities may not wait for a user to click “Confirm”. Thus, MetaMask mode might be most useful for less time-critical arbitrages or for users who want to observe before executing. We will implement it regardless, as it’s a good UX feature, but we educate the user that the fastest opportunities might require pre-authorized auto mode.

- **Private Key (Automated Bot) Integration:** For fully automated operation, the user can provide a private key (or a mnemonic) to the backend, which the MEV bot will use to sign transactions directly. On Replit, we will use **environment variables** or Replit’s Secrets to store this key so it isn’t exposed in the code. Using the web3 library on the backend (for example, Ethers.js), we instantiate a Wallet with this key and connect it to a JsonRpcProvider for each chain. This allows the bot to sign and send transactions programmatically as soon as an opportunity is detected, without user intervention. The advantage is speed; however, managing a raw private key means we must implement strong security practices:

- The key is never logged or transmitted anywhere.
- Only a limited set of operations are done with it (the bot should not, for example, allow arbitrary withdrawal— it only uses the key to perform the MEV strategies defined).
- We might implement an optional safeguard where the user can set spending limits or require a confirmation for trades above a threshold (to prevent, say, a bug from draining funds). For building and sending transactions, we rely on the library’s features. Ethers.js, for instance, automatically handles nonce management per network and can estimate gas. We will double-check nonce synchronization (especially if multiple transactions are sent in quick succession on the same chain, to avoid collisions – a queue per chain might be used). For complex multi-call transactions (like flash loan contract execution), we might pre-deploy a simple smart contract or use an existing one. But given the scope, we try to use direct transactions to existing protocols whenever possible (e.g., calling Uniswap router directly for swaps, calling Aave contract for liquidation).

The bot will use **network RPCs** to broadcast transactions. We will likely use a provider like **Infura** or **Alchemy** for each chain, as they provide reliable endpoints. These services often require their own API keys (Infura project ID, etc.), which we will also store in config. Alternatively, running our own node is possible but not trivial on Replit due to resource constraints. Infura/Alchemy also support WebSocket connections which we use for mempool monitoring and event subscriptions.

#### • Recommended Libraries and Tools:

- For **JavaScript/Node.js**: *Ethers.js* is highly recommended for its ease of use and comprehensive features (signing, contract calls, etc.), and it works smoothly with MetaMask’s provider <sup>21</sup>. *Web3.js* is an alternative, and still widely used <sup>29</sup>, especially if interacting with older guides or tools. Both are capable; Ethers has a cleaner interface for our needs. For interacting with contracts, we’ll include ABI definitions for Uniswap/Sushiswap routers, etc., or use Ethers’

Contract objects with known function signatures. If we use Flashbots, there is an **ethers-provider-flashbots** package that simplifies sending bundles.

- For **Python**: *Web3.py* would be the go-to library <sup>30</sup>. It allows similar functionality in Python. If our backend is Python, we'll rely on *Web3.py* for signing and sending transactions, and perhaps *eth-brownie* or other DeFi libraries for complex interactions. Python has the advantage of libraries like **NetworkX** or **Pandas** if we were to do pathfinding for multi-hop arbitrage or data analysis for scoring, but speed can be a concern for real-time mempool interactions (Node might handle concurrency better for that).
- For **Wallet Connections**: In the frontend, besides MetaMask, we could integrate **WalletConnect** to allow mobile wallets or others to connect by scanning a QR code. This could broaden user access beyond MetaMask extension. There's a WalletConnect web library that could be used with our React app.
- For **Smart Contract interactions**: If we need to deploy or use helper contracts (like a contract to perform multi-step flashloan arbitrage), tools like **Hardhat** or **Truffle** might be used off-line to develop and test them. But within the running app, we'll simply send transactions to those contracts.

- **Transaction Example Flow**: To tie it together, consider an arbitrage example: The bot finds ETH is cheaper on DEX1 than DEX2. It decides to execute. If auto-mode, it constructs two transactions (Swap on DEX1, Swap on DEX2). It estimates gas for each via `estimateGas`. It uses the wallet to sign them. It likely needs to ensure they go in sequence; since both are dependent (you need the output of first to use in second), this could be done by waiting for the first to be mined before sending the second (slower, might miss second-leg profit) or using a contract to do both. A middle approach is using **one transaction with a multi-call** if the DEX supports it (some DEX aggregators or Uniswap v3's router allows multi-hop). Alternatively, use Flashbots to include both in one bundle (so they are mined together or not at all, avoiding leg risk). For now, say we do sequential: the bot sends tx1 with a high gas fee; once it's mined, immediately sends tx2. The bot monitors mempool for any interference (if someone else took the arb in between, tx2 could revert or yield less profit; we handle that by checking state after tx1). In MetaMask/manual mode, the user would be prompted to approve each step – which is likely too slow for profitable arb, but we could design an aggregated transaction that does both trades via a smart contract call so the user approves one transaction. This is an advanced UX (the contract could be a generic "ArbExecutor" contract that given two DEX addresses and a token does the swap sequence).

- **Security Considerations for Execution:**

The app must consider malicious scenarios. For instance, interacting with arbitrary tokens (especially in arbitrage) can be dangerous – the LinkedIn case study <sup>31</sup> describes an MEV bot exploited by a malicious token that had a hidden trap in its contract. To avoid this, our bot will maintain a list of **approved tokens and contracts** (or fetch reputable token lists from sources like CoinGecko or the chain's token registry). If an opportunity involves a token not on the list, the bot might skip it or at least flag it as risky (requiring user approval). We can also simulate trades on a node via `eth_call` (which doesn't change state) to see if something unexpected happens (like the token transfer fails or has strange behavior) before actually sending a transaction. Additionally, by using Flashbots for execution, we can protect ourselves from **sandwiching our own transactions** – if we try a large arbitrage in the public mempool, *other* MEV bots might sandwich or compete, whereas Flashbots bundles are private until mined <sup>22</sup>. We will strive to integrate these best practices into the execution engine to maximize the chances of capturing the value for our user.

## Real-Time Dashboard and Metrics

One of the goals is to provide a rich dashboard that not only shows opportunities but also the performance and behavior of the MEV bot over time. The frontend will display various **metrics and charts** that update in real time:

- **24h MEV Captured:** A running total of profit the bot has made in the last 24 hours (denominated in USD or ETH). This can be shown as a big number and also as a sparkline chart (profit over time). It resets or rolls over after 24h. This metric gives the user an idea of how lucrative the bot's operations are recently.
- **Average Profit per Trade:** This takes all successful MEV trades in the last 24h (or configurable window) and computes the average profit. It's useful to set expectations (e.g., "on average, each arbitrage nets \$5 profit"). It will be displayed next to perhaps a median profit (to see distribution).
- **Win Rate:** The percentage of attempted opportunities that resulted in a successful profit. For example, if the bot attempted 10 executions and 7 made a profit (others failed or broke even/lost gas), the win rate is 70%. We define "win" as net positive profit after gas; a transaction that executed but with no profit or a loss counts as a fail for this stat. This is a crucial measure of the bot's efficiency and can be broken down by category (maybe show win rate for arbitrages vs sandwiches separately, as sandwich attempts might have different risk profiles).
- **Correlation Stats:** To provide deeper insight, we include some correlation analysis. For instance, we can compute the correlation between **Ethereum gas price** and our bot's profitability. If a high gas price correlates with lower profit (likely, since cost is higher and maybe we attempt fewer trades), that's useful info. We can also correlate **market volatility** (perhaps measured by BTC or ETH price swings) with the number of opportunities – often, higher volatility yields more arbitrage and liquidation chances. Another interesting correlation could be between **TVL changes and MEV events** (e.g., did large outflows from a protocol correlate with our liquidation actions?). We will pick a couple of such relationships and show either a scatter plot or a correlation coefficient on the dashboard. For example, a small table might show: "Gas vs Profit: -0.45 (negative correlation)" meaning higher gas hurts profit, or "ETH Volatility vs Opportunities: +0.60 (positive correlation)" indicating more opps in volatile times. These are computed by the backend from data it logs (we'd need to log gas price each time and outcome, etc.).
- **Chain-specific Metrics:** If we support multiple chains, we will show breakdowns per chain. For instance, profits per chain (maybe Ethereum provided \X\$, Polygon \Y\$ in the last day), number of opportunities per chain, etc. This can be a simple bar chart or pie chart.
- **Top Opportunities / Leaderboard:** The dashboard might also highlight the top N profitable trades executed in the past (e.g., "Best Trade: Sandwich on UNI earned \200 at 14:32 UTC") as a bit of a showcase. Conversely, it might list the biggest misses or failures ("Failed arbitrage on JOE, lost \5 gas").

Implementing real-time updates for these metrics will rely on the backend pushing events. Each time the bot completes a trade (success or fail), it can emit an event with the details. The frontend, via WebSocket, receives it and updates the relevant displays (like adding a point to the profit graph or updating the win rate percentage). We'll also have periodic updates for things like correlation which might need a batch of data (maybe computed every hour or on the fly when opened).

From a tooling perspective, we'll likely use **Chart.js** or **ECharts** for plotting graphs in the frontend due to their simplicity and real-time capabilities. For example, Chart.js can update a dataset on the fly causing the graph to animate to the new state. We just need to send the new data points. The UI will be carefully designed so that it remains responsive – perhaps using cards and grids for different metrics, and a modal or separate page for detailed logs if needed.

Finally, all these metrics are not just for show – they also help the user and developers refine the bot. For instance, if the win rate is low, we might decide to raise the confidence threshold or improve the detection accuracy. If the correlation stats show something like “most profit comes from one chain”, the user might allocate more capital there or add another chain with similar characteristics.

## Unified API Layer and Rate Limiting

The **Unified API service** in our backend deserves special mention regarding scalability and reliability. This layer is essentially the gatekeeper for all external data calls. By centralizing API interactions, we gain the following:

- **Rate Limiting and Throttling:** Both CoinGecko and DefiLlama Pro have rate limits on their endpoints (e.g., a certain number of calls per minute depending on plan). Rather than each part of our code independently calling the APIs (which could accidentally exceed limits when combined), the unified service ensures a controlled call rate. We can implement a token bucket or simple queue mechanism. For example, we might allow at most 10 calls per second to CoinGecko; if more requests come, they queue up. This prevents HTTP 429 errors for rate limit and ensures continuous operation. If we do approach limits, the service could also prioritize critical requests (like price queries for an immediate arbitrage) over less critical ones (like an hourly TVL update).
- **Caching Layer:** The unified service caches responses for short durations when appropriate. For instance, CoinGecko's price data might be cached for, say, 5 seconds – in that interval, any component needing price will get it from cache, drastically reducing calls. Similarly, DefiLlama TVL or volume data might only need updating every minute or so. We will fine-tune cache TTLs based on data volatility. The cache can be in-memory (e.g., a simple Python dictionary or Node cache module). For larger or persistent caching, we might use a lightweight Redis (if available on Replit) or a simple file-based cache for things like static lists (coin IDs, protocol list).
- **Unified Response Formatting:** This service also normalizes data formats. CoinGecko and DefiLlama might return data in different structures and units. For instance, CoinGecko prices are in JSON keyed by coin IDs, while DefiLlama might return numbers as strings, etc. The unified API can convert these to a standard format (e.g., all numeric values to JavaScript Number or Python float, timestamps to a common epoch format, etc.). It can also merge data: e.g., combine price from CoinGecko and liquidity from DefiLlama for a given token into one JSON object. This way, the MEV detection logic can call one function `getTokenData(token)` and get everything it needs (price, liquidity, volume, maybe even gas price if we include it) without separately calling multiple APIs.
- **Authentication and Keys:** We store the CoinGecko Pro API key and DefiLlama API key in environment variables on Replit. The unified service attaches these keys to each request (CoinGecko uses an `x-cg-pro-api-key` header or similar, DefiLlama likely similar or a query param for the key). By keeping this in the backend, the frontend or any external party never sees our keys, preventing misuse. If the user needs to provide their own API keys (say they have their

own Pro subscriptions), we could allow configuration of that – but typically one set of keys is fine for the app.

- **Error Handling and Fallbacks:** The service abstracts error handling. If an API call fails (network issue or returns an error), the service can decide to retry or use a fallback. Fallback might mean using a secondary data source (for example, if CoinGecko is down, perhaps switch to an alternate price API like Kaiko or Coinpaprika if we have it integrated). Or use a cached value and log a warning. This makes the overall system more robust – the detection logic can be designed assuming data will come, and the unified layer does its best to fulfill that assumption, or clearly signals an error that can be handled (like throwing a custom exception that the bot knows means “data unavailable, skip this cycle”).
- **Scaling Considerations:** On Replit, we are somewhat limited to a single instance (unless using their new deployment options), but our architecture is inherently scalable. The unified API could be separated into its own microservice if needed, and multiple bot processes could query it. It could also be scaled horizontally by having instances that coordinate through a shared rate-limit counter (though that’s complex). Within one Replit instance, we ensure the unified service is asynchronous (non-blocking IO) so that multiple requests don’t slow each other unnecessarily. For example, using `async/await` in Node or `aiohttp` in Python for concurrent API calls. This is especially useful when, say, we need to fetch data from CoinGecko and DeFiLlama at the same time – we can do both requests in parallel and await both, rather than sequentially, cutting the wait time.
- **Unified API Endpoint (for Frontend):** We may also expose certain data directly to the frontend through our own API endpoints. For example, instead of the frontend calling CoinGecko directly (which it cannot, due to needing the key), the frontend calls our `/api/price?token=XYZ` endpoint. The backend then uses the unified service to get the price and returns it. This way, even ad-hoc requests from the UI (if any) still go through our controlled layer. Generally, though, the frontend will not need to do this often, because we plan to push most needed data proactively.

In summary, the unified API layer is our way to **professionalize** the data handling, much like how an enterprise application would manage external API usage. It ensures that our application can scale in terms of data volume and remain stable even if external services have hiccups or impose limits.

## Multi-Chain and Modular Design

Supporting multiple EVM-compatible blockchains is a core requirement. Our design takes a **modular approach** to multi-chain support, meaning new chains can be added with minimal changes to code. Here’s how we achieve this:

- **Configuration-Driven Chain Setup:** We maintain a configuration (could be a JSON or a Python dict / JS object) that lists each supported chain and relevant parameters. For each chain, we include:
  - *Chain ID and Name:* e.g., Ethereum (ID 1), Polygon (ID 137), Arbitrum (ID 42161), etc.
  - *RPC Endpoint URL:* The connection string to a node on that chain (Infura provides endpoints for some, Alchemy for others; we might use a mix or a service like QuickNode’s multi-chain API). If unavailable, a public RPC or a fallback could be listed too.
  - *DEX Contracts:* A list of the key DEX router addresses and factory addresses on that chain. For example, on Ethereum: Uniswap v2 Router, Sushi Router, Uniswap v3 (if needed with quote),



etc.; on Polygon: QuickSwap, Sushi (again), etc.; on Arbitrum: GMX (for perps?), Uni v3, etc. This list allows the bot to know where to route trades.

- **Token Address Mapping:** For cross-chain comparisons, we need to map tokens across chains (e.g., USDC on Ethereum vs USDC on Polygon – they are different contracts). We keep a mapping of common token symbols to contract addresses per chain. CoinGecko's API often abstracts this if using coin IDs, but for direct contract interactions, we need addresses.
- **Gas Parameters:** Optional overrides for gas limits or gas price strategies per chain. (For instance, on Arbitrum which has very low gas fees and different congestion patterns, we might set a different default strategy than on Ethereum).

This configuration is read by the bot and the unified API service. When the unified service queries DefiLlama or CoinGecko for data, it can use chain-specific identifiers from the config. E.g., DefiLlama might identify chains by name (like "ethereum", "polygon") which we store. The bot uses the RPC endpoint from config to connect a provider for each chain.

- **Chain Abstraction in Code:** We will write functions that are chain-agnostic using the config. For example, instead of hardcoding an Ethereum swap call, we have a function `executeSwap(chain, dex, tokenIn, tokenOut, amount)` that will look up the correct router address from config for the given chain, then use the web3 library to call that router. Similarly, opportunity detection might loop through an array of chains defined in config rather than separate code for each. If certain logic differs (maybe one chain has different dominant DEXs or different thresholds due to gas), we can encapsulate that in the config as well (like a "minProfit" threshold per chain to account for gas differences).
- **Parallel Processing and Thread Safety:** To handle multiple chains at once, our architecture might use multi-threading or asynchronous tasks. For example, the detection engine could spawn a task for each chain's arbitrage scanning so that one slow API call on Polygon doesn't block Ethereum's scanning. We need to ensure that shared resources (like the unified API call counters, or the wallet nonce management if one key is used on all chains) are handled correctly. Typically, nonce and transaction signing are isolated per provider/chain, so that's fine. The unified API service can handle multi-threading by using locks or by an async queue. We will test the system with various chain combinations to ensure adding more chains (scaling out) doesn't degrade performance linearly – ideally, it should handle at least 3-4 chains concurrently given Replit's resource constraints.
- **Extensibility:** When we want to add a new chain (say **Avalanche** or **BSC**), we:
  - Add an entry in the chain config with the chain's name/ID, RPC URL, and list of DEXes etc.
  - Ensure CoinGecko supports price data for that chain's tokens (CoinGecko typically does if the token is listed, as it's by token ID, not chain-specific). If needed, add token addresses for that chain to our token mapping.
  - Ensure DefiLlama has data for that chain (DefiLlama covers many chains for TVL and volume). If the chain is known to DefiLlama, we use the appropriate slug (like "avalanche").
  - Possibly adjust any chain-specific logic (for example, gas calculation formula or known differences in how mempool is accessed – most EVMs are similar though).

Because our detection algorithms mostly rely on data and not chain-specific quirks, they should work out-of-the-box. For instance, arbitrage detection doesn't care if it's Ethereum or Avalanche – if price difference is there and RPC works, it's fine. Sandwich detection needs mempool – that exists on Avalanche too (though fewer bots perhaps). Liquidation detection if we support a new chain means

adding any major lending protocols on that chain to monitor (we might maintain a list of protocol addresses or use DefiLlama's protocol data filtered by chain).

- **Unified View vs Chain-Specific Views:** The dashboard will be able to display either a unified list of opportunities across chains or filter by chain. We will label each opportunity with the chain (and maybe color-code by chain). The user could toggle viewing only Ethereum vs all. This is a UI detail but important for clarity when multi-chain is active.
- **Cross-Chain Opportunities:** While initially out of scope, a modular multi-chain design sets the stage for cross-chain arbitrage (e.g., buy on one chain, bridge, sell on another if price differences are big). This is a complex process involving bridging delays and fees, but as an extension, our architecture could accommodate a *bridging module*. We see DefiLlama also has bridge volume data, which could identify when one chain is cheaper vs another. If implemented, the bot would need to interact with bridge protocols (like Hop or Stargate) to move funds. This is a future consideration, but worth noting that our design's modularity and unified API would help implement such a feature (since we already track multi-chain data and could integrate a bridging API similarly).

In essence, modular multi-chain support means our app is not hard-coded to Ethereum. This not only broadens the scope of opportunities (some MEV opportunities might be easier on smaller chains with fewer competitors), but also makes the system more resilient – if one network is too congested or not profitable, the bot might find value on another.

## Security Considerations

Security is critical in an application that handles private keys, financial value, and interacts with potentially adversarial environments (blockchain mempools). We address security at multiple levels:

- **Key Management:** The user's private key (if provided for automation) is stored securely on the backend. Replit allows storing secrets which are not exposed in code. We will use that mechanism or an environment variable. The key is loaded into memory by the bot and used for signing. We will *never log this key* or transmit it. The front-end will not have access to it. If using MetaMask, the private key security is delegated to MetaMask <sup>20</sup>, which is a trusted wallet; our app just sends transaction requests to it. We will also advise users to use a dedicated wallet with limited funds for this bot, to limit risk (so even if compromised, it's not their life savings).
- **API Keys Security:** Similar to above, the API keys for CoinGecko and DefiLlama are stored in the backend. The frontend calls our backend for data, never directly the third-party APIs, so the keys stay secret. We also ensure our repository (if the Replit is shared or exported) does not contain those keys.
- **Secure Communication:** The app will be served over HTTPS (Replit sites typically have HTTPS by default). This ensures that the data between the user's browser and our server (including any wallet addresses or triggered transactions) is encrypted in transit. If we integrate MetaMask, the communications go through the browser extension which is also secure. We should be cautious about any mixed content or external resources; e.g., if we embed images or fetch scripts, ensure they come from HTTPS sources to avoid man-in-the-middle injection.
- **Pre-trade Validation:** As mentioned earlier, the bot will simulate or validate opportunities before execution to avoid traps. This includes:

- Checking token contract code or using known token lists to avoid **malicious tokens** (honeypots that block selling, or have hidden transfer fees, etc.). We might integrate with a service or have a list of known scams to skip.
- Using Flashbots for suspicious transactions to avoid exposing them in mempool – if we suspect an opportunity might be a bait (like the case where an arbitrage was intentionally made to lure bots <sup>31</sup>), by submitting via Flashbots we either get it executed atomically or not at all, potentially avoiding the trap of on-chain trickery (though a smart contract trap would still execute in Flashbots if we call it, so the primary defense is not calling unverified contracts).
- Ensuring that for any liquidation, the target protocol's contracts are authentic (we will rely on known addresses or DefiLlama's data which likely uses verified protocol addresses). Many phishing contracts try to mimic real ones; we'll only interact with addresses from official sources.
- **Transaction Ordering and Front-Run Protection:** When our bot finds an opportunity and is about to act, we have to consider that *we ourselves* might get front-run by others. To mitigate this:
  - We often use **Flashbots** for critical transactions (especially sandwiches, as discussed, and possibly large arbitrages) to hide our intent from the public mempool <sup>22</sup>. This prevents others from seeing our transaction beforehand.
  - If not using Flashbots, we set a high priority fee to increase the chance our transaction is mined quickly. We also monitor the mempool; if we see another bot submitted a similar trade with higher gas, we might choose to cancel ours (if possible) to save gas.
  - The bot might include small bribes in its Flashbots bundles to incentivize inclusion (Flashbots bundles include a fee payment to the miner outside of the block rewards).
- **Error and Revert Handling:** The bot will handle transaction failures gracefully. A reverted transaction (e.g., due to slippage tolerance or someone else executing the arb first) is an expected risk. We will set sane parameters (like slippage limits on DEX trades) to avoid overpaying. If a transaction reverts, our bot logs it, does not continually retry blindly (which could waste gas repeatedly), and possibly raises the opportunity's risk level. If using Flashbots, a failed bundle simply won't be mined, so no gas is lost in that case (another benefit of Flashbots: you only pay if it succeeds, in general).
- **Secure Development Practices:** We will use reputable libraries and keep them updated (CoinGecko and DefiLlama usage doesn't require third-party SDKs, just HTTPS calls which are fine). If any user input is taken (like entering a private key or adjusting a parameter), we validate and sanitize it. The front-end and back-end are under our control entirely, so XSS or injection risks are low, but we still will not directly inject any untrusted data into pages without encoding. For example, if we ever displayed data from an API on the page, we ensure it's properly handled (though most data here are numeric or known strings).
- **Resource Usage and Denial-of-Service:** The unified API layer helps prevent external APIs from DoS'ing us by rate limiting, but we also consider our own app's exposure. If the app is public, someone could spam our endpoints (like calling `/api/execute` rapidly). We may implement basic rate limiting on sensitive endpoints and perhaps an authentication token for the execute endpoint (so only the legitimate user triggers it from the UI). Replit applications can have usage limits, and if our app gets overloaded (maybe by an infinite loop or external attack), it could crash or be rate-limited by Replit. We will test for memory leaks or unbounded queues (for

example, ensure if an API is slow, we don't just accumulate 1000 pending requests). Scalability has been addressed in architecture; here it overlaps with security in terms of availability.

- **Replit Deployment Specifics:** Running on Replit has some implications. We should use Replit's secrets for API keys. We should be mindful that Replit might restart our server occasionally or after code edits – we should persist important state if any (maybe not much state beyond logs; we can persist logs or metrics in a small database or even a JSON file on the Replit persistent storage `/mnt/data`). We will also set up proper environment variables for configuration (like RPC URLs, since those might contain keys from Infura). Replit's environment is containerized, so it's relatively secure from outside interference aside from our exposed web endpoints. We will update dependencies to avoid any known vulnerabilities in packages.
- **User Education and Control:** Finally, we incorporate some features to let the user remain in control for safety. For example, a "kill switch" to pause the bot (maybe a button on the UI that flips an `active` flag off so the bot stops taking new actions – useful if things seem to go wrong or market conditions suddenly change drastically). Also, transparency: the user can see what the bot is about to do. For manual mode, this is obvious (they confirm each tx). For auto, we still log the details in the UI (like "Attempting sandwich on swap X, gas price Y, expecting profit Z"). This way the user isn't blindsided by the bot's actions on their behalf.

By addressing security at these multiple layers, we aim to protect both the user's funds and the integrity of the system. The blockchain space is adversarial, so continuous monitoring and updates will be necessary as new threats (or MEV techniques) emerge, but the framework put forth is robust and secure by design.

## Conclusion

In this report, we designed a comprehensive framework for an MEV opportunity identification and execution system suitable for deployment on Replit. The proposed architecture integrates data from CoinGecko Pro and DefiLlama Pro to detect a variety of MEV strategies (sandwich attacks, arbitrage trades, and liquidation sniping) across multiple EVM chains in real time. We detailed how the **frontend dashboard** provides users with actionable insights and controls, while the **backend** (augmented by a unified API layer) handles data aggregation, opportunity detection, and transaction execution via a secure **MEV bot engine**.

We identified key API endpoints from CoinGecko (for live pricing and market info) and DefiLlama (for DeFi analytics like TVL, liquidity, and user flows) and explained how each contributes to the system's functionality – from fetching prices to calculating profits. The framework emphasizes modularity, allowing easy expansion to new chains and strategies, and scalability, using techniques like unified request handling and caching to manage external API usage.

Crucially, we incorporated robust **security considerations**: leveraging MetaMask for secure key management <sup>20</sup>, using private keys carefully with trusted libraries <sup>21</sup>, protecting transactions from competitors (e.g., via Flashbots bundles <sup>22</sup>), and avoiding pitfalls like malicious contracts <sup>31</sup>. Real-time monitoring and user-defined parameters add additional layers of safety and flexibility.

On Replit, this application can be deployed as a persistent web service with scheduled tasks. Replit's always-on capability (with proper plan) means the bot can run continuously, and the web UI will be available for the user to monitor and control the operations. The design is mindful of Replit's

environment, using environment variables for secrets and efficient single-instance operation for now, with potential to split components if needed in the future.

By combining the rich datasets of CoinGecko and DefiLlama with a tailored MEV searcher bot, this framework offers a powerful tool for users to capture value in DeFi markets. It is both **fully functional** – covering detection through execution – and **scalable**, capable of handling multiple chains and high-throughput data streams. With a clear separation of concerns (frontend vs backend vs data layer vs bot) and thorough planning of each part, the implementation can proceed with confidence. This solution not only meets the requirements but also provides a blueprint for a professional-grade MEV platform that can adapt to the evolving DeFi landscape and remain secure and effective over time.

---

1 2 3 4 7 9 22 31 Understanding MEV Bots: Impact, Incidents, and Deep Dive into the Mempool

<https://www.linkedin.com/pulse/understanding-mev-bots-impact-incidents-deep-dive-mempool-singh-pktof>

5 10 16 23 24 25 26 Endpoint Overview

<https://docs.coingecko.com/reference/endpoint-overview>

6 8 11 12 13 14 15 17 18 27 28 defillama-api-spec (2).json

<file:///file-UagWZdQTyc5RhVgbHjEDYw>

19 20 Web3.js and MetaMask Integration: A Comprehensive Guide | by asierr.dev | Medium

<https://medium.com/@asierr/web3-js-and-metamask-integration-a-comprehensive-guide-12b8f7fd97e3>

21 29 30 Web3 libraries | MetaMask developer documentation

<https://docs.metamask.io/services/concepts/web3-libraries/>